

## LAB-1

# Calculate Prediction Confidence Intervals for a Linear Regression Model

By

Dr. Sumit Kumar Singh

Associate Professor

SR University





# WHAT IS REGRESSION?

- Regression helps us understand how one variable changes when another changes.

Example:

Study hours → Exam marks

House size → House price

- Regression draws a best-fitting line through the data to model this relationship.

## EXAMPLE: STUDY HOURS VS MARKS

Study Hours (X)	Marks (Y)
1	45
2	50
3	55
4	60
5	65

Observation: As study hours increase, marks increase, the pattern looks like a straight line.

# WHAT IS LINEAR REGRESSION?

➤ Linear Regression assumes a straight-line relationship:

$$Y = \beta_0 + \beta_1 X + \varepsilon$$

Where:

Y – Dependent variable (output)

X – Independent variable (input)

$\beta_0$  – Intercept (value of Y when X=0)

$\beta_1$  – Slope (change in Y for one unit change in X)

$\varepsilon$  – Random error

**Multiple Linear Regression (MLR)** extends this to **two or more independent variables**:

$$Y = a + b_1X_1 + b_2X_2 + b_3X_3 + \cdots + b_nX_n + \epsilon$$

**What is  $R^2$  (R-squared)?**

- **$R^2$  (Coefficient of Determination)** measures **how well the regression line fits the data**.
- It tells us: “How much of the variation in Y (Marks) is explained by X (Study Hours).”

$R^2=1 \rightarrow$  Perfect linear relationship

$R^2=0 \rightarrow$  No relationship

$R^2=0.8 \rightarrow$  80% of variation explained by model

## FORMULA

$$R^2 = 1 - \frac{SS_{res}}{SS_{tot}}$$

$$SS_{res} = \sum (Y - \hat{Y})^2$$

$$SS_{tot} = \sum (Y - \bar{Y})^2$$

$SS_{tot}$  : How much the actual data vary around their mean (total variation).

$SS_{res}$  : How much error remains after fitting the regression model (unexplained variation).

# WHAT IS A CONFIDENCE INTERVAL?

- A confidence interval (CI) gives a range within which we expect the true value to fall.

Example:

If predicted marks = 70,  $CI = [68, 72]$ ,  
then we are 95% confident the true mean lies between 68 and 72.

# CONFIDENCE INTERVAL FORMULA

$$\hat{Y} \pm t_{(\alpha/2, n-2)} \times SE(\hat{Y})$$

Where:

$\hat{Y}$  – Predicted value

$SE(\hat{Y})$  – Standard error of prediction

$t_{(\alpha/2, n-2)}$  – t-critical value


For a 95% confidence interval,  
we're 95% confident the true value is inside the  
range.

so  $\alpha = 1 - 0.95 = 0.05$ .

So,  $\alpha = 0.05$  represents the remaining 5% risk

Common choice:  $\alpha = 0.05 \rightarrow 95\%$  confidence level.




$$\hat{Y} \pm t_{(\alpha/2, n-2)} \times \text{SE}(\hat{Y})$$

Why  $\alpha/2$  ?

➤ When we make a **two-sided confidence interval** (both upper and lower limits), we divide that 5% risk equally into **two tails** of the t-distribution:

2.5% in the **left tail**

2.5% in the **right tail**

So, we use  $\alpha/2 = 0.025$  for each side

Why  $n - 2$ ?

if there are  $n$  total observations, and we estimated 2 parameters:

$$\text{Degrees of Freedom} = n - 2$$

The **t-distribution** depends on the number of degrees of freedom.

Fewer data points (small  $n$ )  $\rightarrow$  wider tails  $\rightarrow$  larger t-critical value (more uncertainty).

So,  $t_{\alpha/2, n-2}$  automatically adjusts based on sample size.



# STEPS IN PYTHON

- Import libraries and dataset
- Build the LR model
- Make prediction and confidence intervals
- Visualize Regression line with Confidence band

# STEP 1: IMPORT LIBRARIES AND DATASET

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import statsmodels.api as sm
```

**Python library for statistical analysis, data exploration, and regression modeling.**

```
data = {
    'Hours': [1, 2, 3, 4, 5, 6, 7, 8],
    'Marks': [45, 48, 52, 56, 61, 65, 68, 72]
}
df = pd.DataFrame(data)
print(df.head())
```

## STEP 2: FIT LINEAR REGRESSION MODEL

```
X = sm.add_constant(df['Hours'])    # adds intercept term  
y = df['Marks']
```

```
model = sm.OLS(y, X).fit()  
print(model.summary())
```

sm.OLS() = Ordinary Least Squares regression

model.summary() shows slope, intercept, and model accuracy ( $R^2$ )

## STEP 3: MAKE PREDICTIONS & COMPUTE CONFIDENCE INTERVALS

```
pred = model.get_prediction(X)
conf_int = pred.conf_int(alpha=0.05)
predicted = model.predict(X)
```

```
result = df.copy()
result['Predicted'] = predicted
result['Lower_CI'] = conf_int[:,0]
result['Upper_CI'] = conf_int[:,1]
```

```
print(result.head())
```

➤ `model.get_prediction(X)` → This line asks the fitted regression model to generate predictions and statistical details (like confidence intervals) for each input in X.

It doesn't just give predicted values, it returns a **PredictionResults object**, which contains:

Mean predicted values  $\hat{Y}$

Standard errors of predictions

Confidence intervals

➤ `conf_int = pred.conf_int(alpha=0.05)` → This extracts the **confidence intervals** (CIs) from the prediction results.


$\alpha=0.05$  means a 95% confidence level (since  $1-\alpha=0.95$ ).

The result is an array with two columns:

Column 0 → Lower bound of the 95% CI

Column 1 → Upper bound of the 95% CI

➤ `predicted = model.predict(X)` → This line generates the **predicted values**  $\hat{Y}_i$  from your regression model, the values that lie exactly on your best-fit line.



Line	Command	What It Does
1	<code>model.get_prediction(X)</code>	Generates predicted means + errors + intervals
2	<code>pred.conf_int(alpha=0.05)</code>	Extracts lower & upper 95% confidence bounds
3	<code>model.predict(X)</code>	Calculates predicted (fitted) values $\hat{Y}$
4	Add new columns to df	Combines original data with predicted + CI
5	<code>print(result)</code>	Displays the result summary table

## STEP 4: VISUALIZE REGRESSION LINE WITH CONFIDENCE BAND

```
plt.figure(figsize=(8,5))
plt.scatter(df['Hours'], df['Marks'], color='blue', label='Actual')
plt.plot(df['Hours'], predicted, color='red', label='Predicted line')
plt.fill_between(df['Hours'], result['Lower_CI'], result['Upper_CI'], color='gray', alpha=0.3,
label='95% Confidence Interval')
plt.xlabel('Study Hours')
plt.ylabel('Marks')
plt.title('Linear Regression with 95% Confidence Interval')
plt.legend()
plt.show()
```





# INTERPRETING THE GRAPH

Red line: Model's predicted values.

Gray band: 95% confidence interval around predictions.

Narrow band → more certainty.

Wide band → less certainty.